

# Azure SQL MI PowerBI

March 2019

## Contents

Introduction.....	1
SQL Managed Instance Environment.....	2
Exercise 0: Using Pre-configured Environment .....	2
Overview .....	2
Exercise 1: Login to Azure Portal .....	<b>Error! Bookmark not defined.</b>
Overview .....	<b>Error! Bookmark not defined.</b>

## Introduction

In this lab, you will learn key concepts in data integration centered on orchestrating and operationalizing data movement and data transformation activities, as well as cloud-based SSIS package execution via Azure Data Factory's Integration Runtime feature. You will establish connections with both on-premises SQL Server Databases, as well as Managed Instances in the cloud, to allow the creation of data pipelines that run from your sources, all the way to your destination. Then, you will lift-and-shift native SSIS packages onto the cloud to allow for more seamless control and monitoring during executions. Finally, you'll complete the journey of your data by establishing a connection to a live Power BI Desktop report, where you will be able to visualize directly the various components and insights contained in your data.

### Understanding Integration Runtime in Azure Data Factory

**Integration Runtime (IR)** *is the **compute infrastructure** used by Azure Data Factory to provide the following data integration capabilities across different network environments:*

- **Data Movement:** Moving data between data stores in public and private networks.
- **Activity Dispatch:** Monitoring transformation activities running on a variety of compute services.
- **SSIS Package Execution:** Natively execute SQL Server Integration Services (SSIS) packages in a managed Azure compute environment.

**Azure Data Factory** offers three types of Integration Runtimes that serve specific data integration capabilities; Azure, Self-hosted, and **Azure-SSIS**. The last type supports the previously mentioned **SSIS Package Execution**, which is what we'll be using in this section of the lab.

# SQL Managed Instance – Power BI

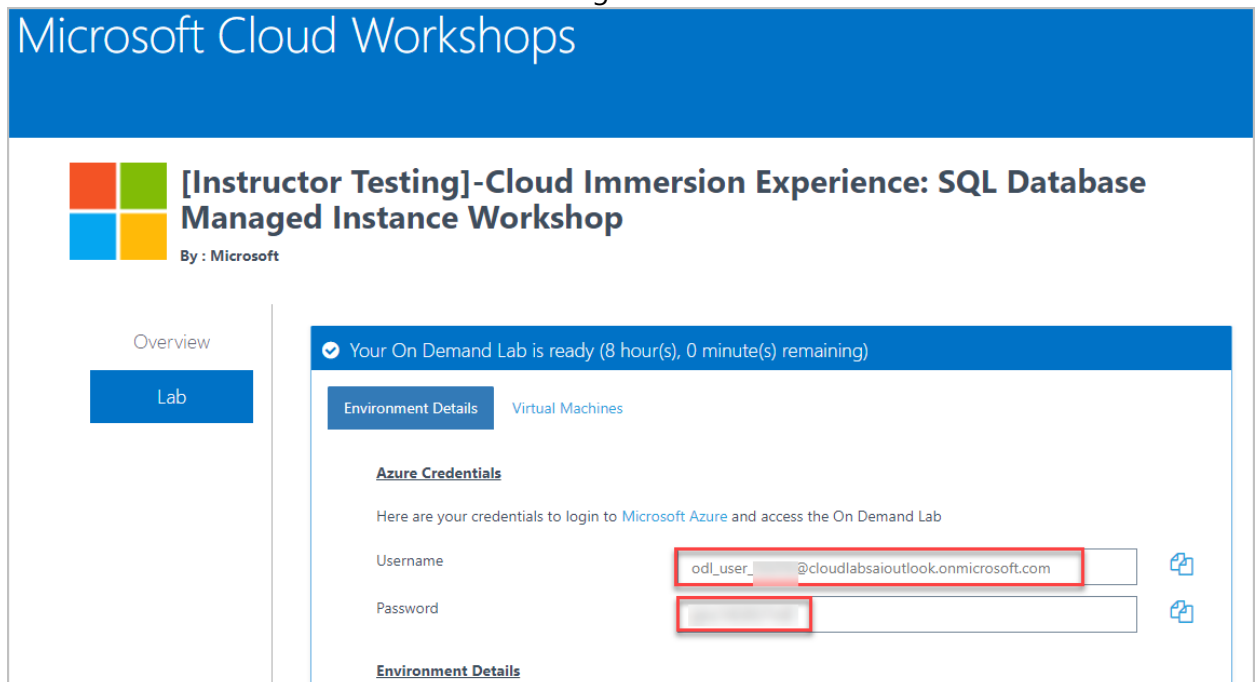
## Exercise 0: Provisioning an Azure-SSIS Integration Runtime in Azure Data Factory

### Overview

This section will provide you with the ability to use the Azure Portal to create the foundation needed for the Integration Runtime's infrastructure to effectively host and manage the executions of SSIS packages residing virtually anywhere in the world.


### Task 1: Configure Integration Runtime Setup

1. Launch a browser and navigate to <https://portal.azure.com>. Once prompted, login with the Azure Credentials from the Lab Details Page.




2. In the Stay signed in? pop-up window, click No
3. In the Welcome to Microsoft Azure pop-up window, click Maybe Later  
Note: If you receive a pop-up for Azure Advisor, click the X in the top right corner of the pop-up to close it.
4. You will be directed to the dashboard.
5. From the left side of the Page, select **Resource Groups**

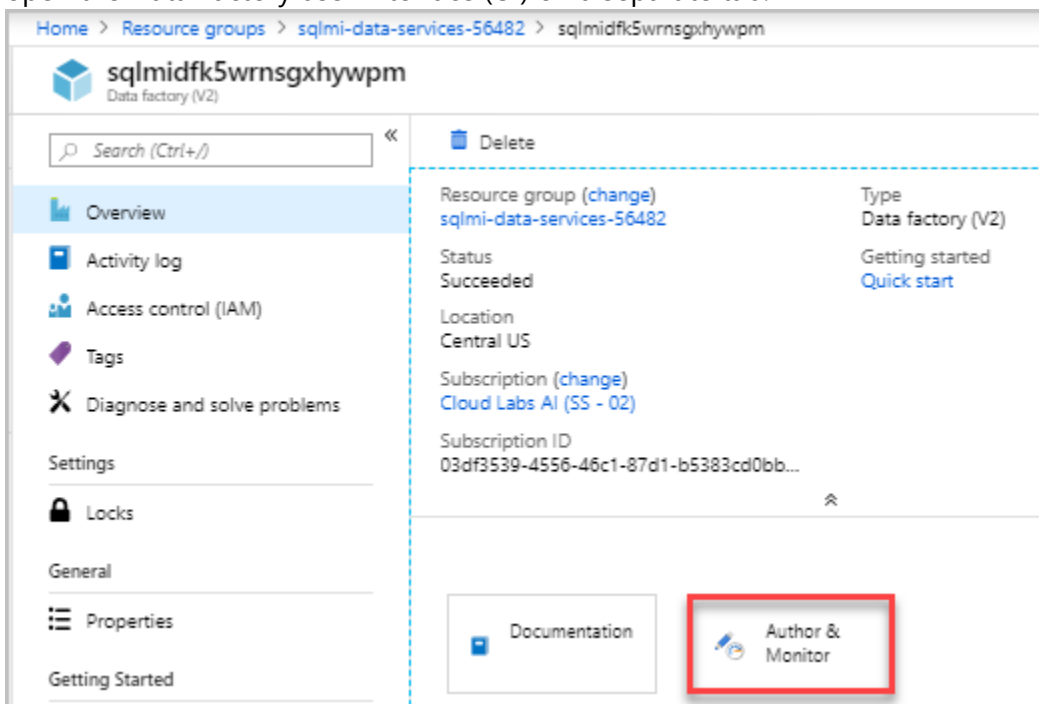
6. Note that you will have access to three Resource groups:




7. Select resource group starting with **sqlmi-data-services-**  
8. From the overview blade of the Resource group, **select** the resource with type **Data Factory (V2)** from the resources:



9. From the overview blade of the Data Factory, navigate to the Author & Monitor tile to open the Data Factory user Interface (UI) on a separate tab.



10. On the **Let's get started** page, click the **Configure SSIS Integration Runtime** tile



11. On the **General Settings** form of Integration Runtime Setup page, enter the following information:
  - a. **Name:** integrationRuntime1
  - b. **Description:** Leave blank
  - c. **Location:** Select **location of the resource group** from the dropdown

**NOTE:** The location does not need to be that of your data factory's, but it should be the same as your Azure SQL Database/Managed Instance server's where your **SSISDB** is to be hosted, allowing easy access without incurring excessive traffic.

- d. **Node Size:** Select **Standard\_E8\_v3** from the dropdown \*This node's high memory-to-CPU ratio results in superior performance when working with **large database servers** - which this lab will. Another case where selecting a large node size would be beneficial is when you anticipate running **compute/memory-intensive SSIS packages**
- e. **Node Number:** Select **1** by using the slider. *The number of nodes you select determines the cluster size of your IR, which impacts performance when running parallel SSIS packages. **Select a large cluster if you plan on running many packages in parallel***
- f. **Edition/License:** Select **Enterprise** from the dropdown, as it will allow the use of advanced/premium features on your integration runtime, showcasing various features during this lab
- g. **Save Money:** Click **No**

## Integration Runtime Setup

### General Settings

Name \*

integrationRuntime1

Description

Type

Azure-SSIS

Location \*

Central US

Node Size \*

Standard\_E8\_v3 (8 Core(s), 65536 MB)

Node Number \*

1

Edition/License \*

Enterprise

Save Money

Save with a license you already own. Already have a SQL Server license?

Yes

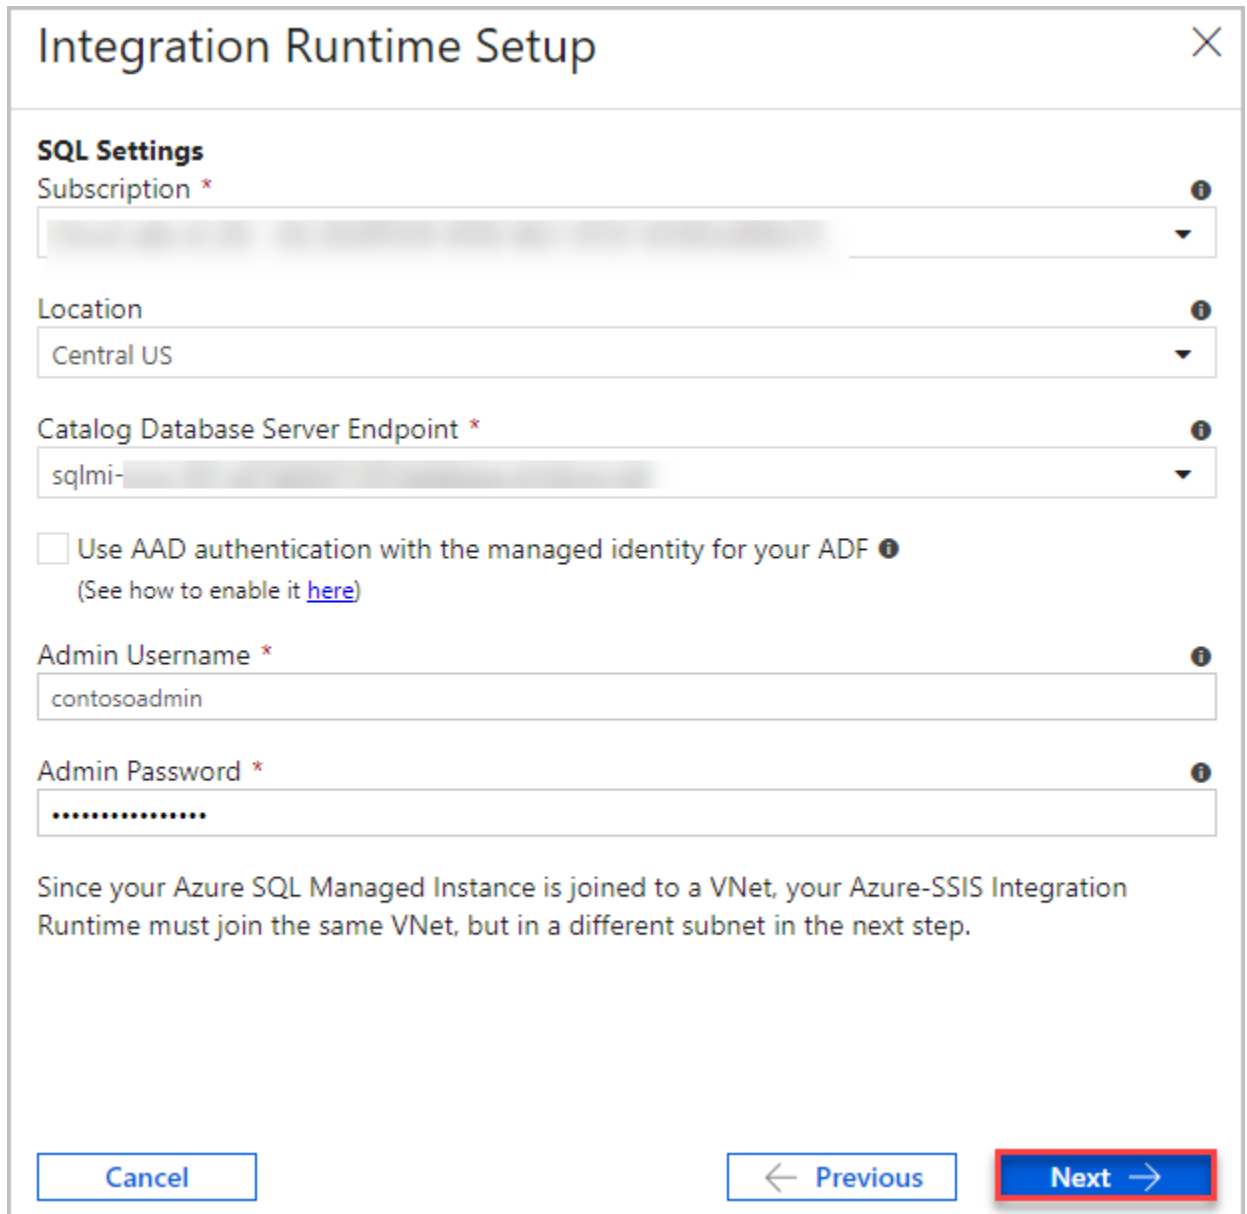
No

Cancel

Next →

12. Click **Next** to bring up the **SQL Settings** portion of Integration Runtime Setup page, and provide the following information:
13. **Subscription:** Select your default subscription
14. **Location:** Select your resource group location from the dropdown, as it is recommended that you select the same location of your integration runtime
15. **Catalog Database Server Endpoint:** Select the available Database Server Endpoint from the Dropdown.
16. **Do not** check **Use AAD authentication with your ADF MSI**
17. **Admin Username:** contosoadmin

18. **Admin Password:** IAE5fAijit0w^rDM



The image shows a screenshot of the 'Integration Runtime Setup' dialog box. It has a title bar with a close button (X). The main content area is divided into sections. The first section is 'SQL Settings', which includes three dropdown menus: 'Subscription' (blurred), 'Location' (set to 'Central US'), and 'Catalog Database Server Endpoint' (set to 'sqlmi-'). Below these is a checkbox labeled 'Use AAD authentication with the managed identity for your ADF' with an information icon and a link '(See how to enable it [here](#))'. The next section is 'Admin Username' with a text box containing 'contosoadmin'. Below that is 'Admin Password' with a masked text box showing dots. At the bottom, there is a paragraph of text: 'Since your Azure SQL Managed Instance is joined to a VNet, your Azure-SSIS Integration Runtime must join the same VNet, but in a different subnet in the next step.' At the very bottom, there are three buttons: 'Cancel', 'Previous' (with a left arrow), and 'Next' (with a right arrow and a red border).

**Integration Runtime Setup**

**SQL Settings**

Subscription \*

Location

Central US

Catalog Database Server Endpoint \*

sqlmi-

☐ Use AAD authentication with the managed identity for your ADF ⓘ  
(See how to enable it [here](#))

Admin Username \*

contosoadmin

Admin Password \*

.....

Since your Azure SQL Managed Instance is joined to a VNet, your Azure-SSIS Integration Runtime must join the same VNet, but in a different subnet in the next step.

Cancel Previous Next

19. **Maximum Parallel Executions Per Node:** select **1** from the dropdown box

**NOTE:** the maximum number of packages to execute concurrently per node in your integration runtime cluster. Only supported package numbers are displayed. Select a low number, if you want to use more than one cores to run a single large/heavy-weight package that is compute/memory -intensive. Select a high number, if you want to run one or more small/light-weight packages in a single core.

20. **Custom Setup Container SAS URI:** Leave blank

21. **Select a VNet...:** Leave checked



**NOTE:** Selecting whether you want to join your integration runtime to a virtual network: You should check it if you use Azure SQL Database with virtual network service endpoints/Managed Instance to host SSISDB or require access to on-premises data.

22. Select the subscription from the dropdown.

## Integration Runtime Setup

### Advanced Settings

Maximum Parallel Executions Per Node \*

1

Custom Setup Container SAS URI

☒ Select a VNet for your Azure-SSIS Integration Runtime to join and allow ADF to create certain network resources  
(See more info [here](#))

Subscription \*

Cloud Labs AI (SS - 02) (03df3539-4556-46c1-87d1-b5383cd0bb27)

Location \*

Central US

Type \*

Azure Resource Manager Virtual Network

VNet Name \*

SQLMI-VNET-Shared

Subnet Name \*

sqlmi-dms-subnet

Cancel

VNet Validation

← Previous

Next

23. Click **VNet Validation** and if successful, click **Next**

✓ Validation successful!

Cancel

VNet Validation

24. Click **Finish** to start the creation of your integration runtime

## Integration Runtime Setup

### Summary

Your Azure-SSIS Integration Runtime (IR) is created with the following settings:

#### Azure Data Factory Settings

- **Subscription:** 03df3539-4556-46c1-87d1-b5383cd0bb27
- **Resource Group:** sqlmi-data-services-56482
- **Name:** sqlmidfk5wrnsgxhywpm
- **Location:** centralus

#### General Settings

- **Name:** integrationRuntime1
- **Location:** Central US
- **Node Size:** Standard\_E8\_v3
- **Node Number:** 1
- **Edition:** Enterprise
- **Azure Hybrid Benefit:** LicenseIncluded

#### SQL Settings

- **Catalog Database Server Endpoint:** sqlmi-scus-001.a07ab820155f.database.windows.net
- **Catalog Database Server Location:** centralus

#### Advanced Settings

- **Maximum Parallel Executions Per Node:** 1
- **VNet Name:** SQLMI-VNET-Shared
- **Subnet Name:** sqlmi-dms-subnet

If you want to change any of the above settings, click **Previous** to do so.







Cancel

← Previous

Finish

25. On the **Connections** tab, switch to **Integration Runtimes** if needed. Select **Refresh** to refresh the status

The screenshot shows the Microsoft Azure Data Factory interface. The top navigation bar includes 'Data Factory', 'Publish All', 'Validate All', 'Refresh', and 'Discard All' buttons. The left sidebar shows 'Factory Resources' with 'Pipelines' and 'Datasets' counts. The main area shows the 'Connections' tab with the 'Integration Runtimes' sub-tab selected. The 'Integration Runtimes' sub-tab is highlighted with a red box. Below the sub-tab, there is a '+ New' button and a 'Refresh' button. A table lists the integration runtimes:

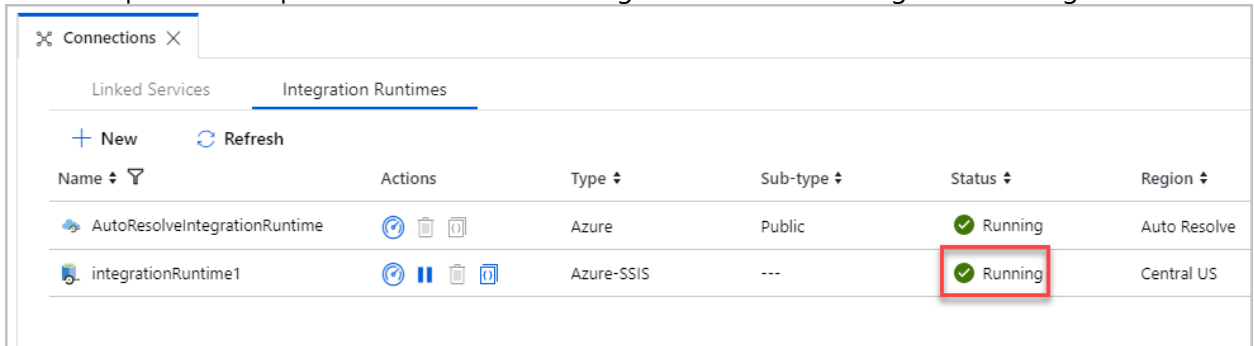
Name	Actions	Type	Sub-type	Status	Region
AutoResolveIntegrationRuntime	  	Azure	Public	Running	Auto Resolve
integrationRuntime1	  	Azure-SSIS	---	Starting	Central US








**NOTE:** This process takes approximately **20 to 30 minutes** to complete due to the Azure Feature Pack for SSIS and the Access Redistributable installations. This is taking place while the Data Factory service connects to your Azure SQL Database server to prepare the **SSIS Catalog** (SSISDB database).

**NOTE:** Use the links in the **Actions** column to stop/start, edit, or delete the integration runtime. Use the last link to view JSON code for the integration runtime. The edit and delete buttons are enabled only when the IR is stopped.

26. Please continue with the lab while the integration runtime installs

27. Once the process completes the status of the integrationRuntime1 changes to Running.



Connections					
Linked Services		Integration Runtimes			
+ New Refresh					
Name	Actions	Type	Sub-type	Status	Region
AutoResolveIntegrationRuntime	  	Azure	Public	Running	Auto Resolve
integrationRuntime1	   	Azure-SSIS	---	Running	Central US

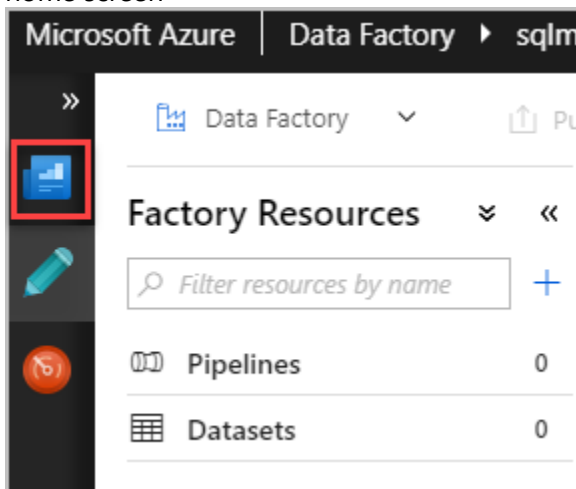
## Exercise 2: Use the Copy Data Tool to Create and Run Data Pipeline

### Overview

I

#### Task 1:

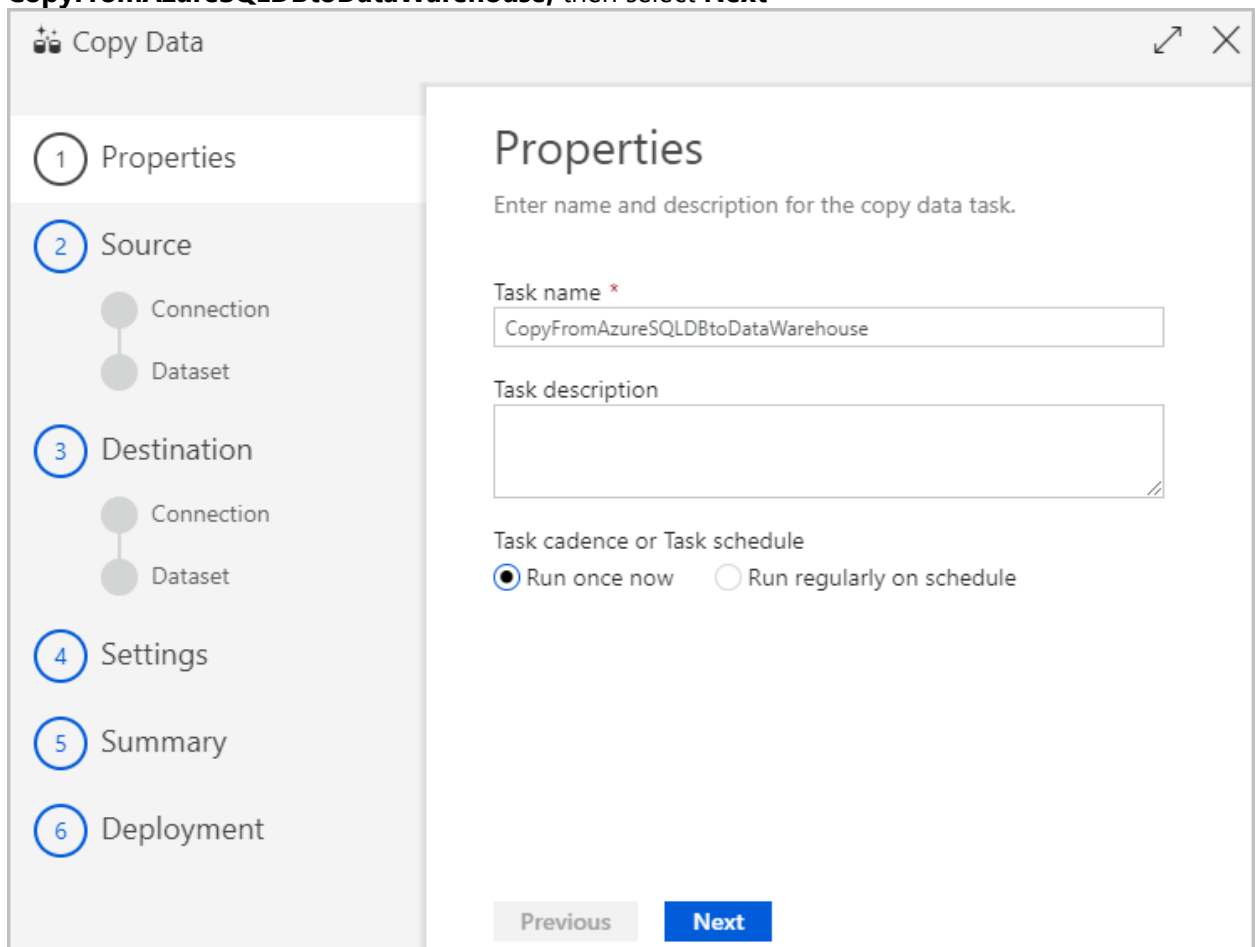
1. On the left-hand column of the screen, click the blue **Data Factory** icon to return to the home screen



2. On the **Let's get started** page, select the **Copy Data** tile to launch the **Copy Data tool**



3. On the **Properties** page, under **Task name**, enter **CopyFromAzureSQLDBtoDataWarehouse**, then select **Next**

A screenshot of the 'Copy Data' tool's 'Properties' configuration window. The window has a title bar 'Copy Data' with standard window controls. On the left is a vertical sidebar with six steps: 1 Properties, 2 Source, 3 Destination, 4 Settings, 5 Summary, and 6 Deployment. Step 2 'Source' is currently selected. Within 'Source', 'Connection' and 'Dataset' are listed with radio buttons; 'Connection' is selected. Steps 3 through 6 also show 'Connection' and 'Dataset' options. The main area is titled 'Properties' and contains instructions: 'Enter name and description for the copy data task.' It includes a 'Task name' field with a red asterisk, containing the text 'CopyFromAzureSQLDBtoDataWarehouse'. Below it is a 'Task description' text area. At the bottom, there are radio buttons for 'Task cadence or Task schedule': 'Run once now' (selected) and 'Run regularly on schedule'. At the very bottom are 'Previous' and 'Next' buttons, with 'Next' being highlighted in blue.


4. On the **Source data store** page, click **+ Create new connection**

**Source data store**

Specify the source data store for the copy task. You can use an existing data store connection or specify a new data store.

All   Azure   Database   File   Generic protocol   NoSQL   Services and apps

All   Filter by name   **+ Create new connection**












No connection to display.

Try changing your filters if you don't see what you're looking for.

Previous   Next

5. Select **Azure SQL Database Managed Instance** from the gallery, and then select **Continue**

**New Linked Service** ✕

 Azure Database for PostgreSQL	 Azure File Storage	 Azure SQL Data Warehouse
 Azure SQL Database	 Azure SQL Database Managed Instance	 Azure Table Storage
 cassandra		 cassandra

Cancel   Continue

6. On the New Linked Service page, enter the following information:
  - a. **Name:** Enter **CopySourceService**
  - b. **Description:** Leave blank
  - c. **Connection via Integration Runtime:** Select **AutoResolveIntegrationRuntime**
  - d. **Connection String or Azure Key Vault:** Select **Connection String**
  - e. **Account Selection Method:** Select **From Azure subscription**
  - f. **Azure subscription:** select **default subscription**
  - g. **Server name:** Provide SQL MI Server
  - h. **Database Name:** Provide **the database name of the database you migrated**
  - i. **Authentication Type:** Select **SQL Authentication**
  - j. **User name:** contosoadmin
  - k. **Password:** **IAE5fAijit0w^rDM**
  - l. **Click on Test**
  - m. Click **Finish**
7. Select the newly created linked service as source, then click **Next**
8. On the **Select tables from which to copy the data or use a custom query** page, select the Table **[Purchasing].[PurchaseOrders]**, then click **next**
9. On the **Destination data store** page, click + **Create new connection** to add a connection
10. Select **Azure SQL Data Warehouse** from the gallery, and then select **Continue**
11. On the **New Linked Service** page, enter the following information:
12. **Name:** Enter **CopyDestinationService**
13. **Description:** Leave blank
14. **Connection via Integration Runtime:** Select **AutoResolveIntegrationRuntime**
15. **Connection String or Azure Key Vault:** Select **Connection String**
16. **Account Selection Method:** Select **From Azure subscription**
17. **Azure subscription:** select
18. **Server name:** Provide SQL MI Server Name
19. **Database Name:** Select **DataWarehouse-**
20. **Authentication Type:** Select **SQL Authentication**
21. **User name:** Enter **LabUser**
22. **Password:** Enter
23. Click **Finish**
24. Select the newly created linked service as sink, then click **Next**
25. On the **Table mapping** page, select **Next**
26. On the **Column mapping** page, select **Next**
27. On the **Settings** page, Uncheck **Enable Staging**, and **Allow polybase**
28. Click **Next**
29. On the **Summary** page, review the settings, and then select **Next**
30. On the **Deployment** page, select **Monitor** to monitor the pipeline (task)

**NOTE:** Notice that the Monitor tab on the left is automatically selected. The Actions column includes links to view activity run details and to rerun the pipeline

## Exercise 3: Deploy and Run SSIS Packages from Azure Data Factory SSIS IR

### Overview

I

#### Task 1:

Verify the Integration Runtime you created earlier in the lab has completed.

1. On the monitor page, click the **Integration Runtimes** tab near the top of the page
2. Verify that **MyFirstIntegrationRuntime** has a Status of **Running**. If its status is **Starting** then wait for the Runtime to complete before continuing.

**NOTE:** Click the refresh button to update the status

**NOTE:** This process takes approximately **20 to 30 minutes**

### Connect to SSISDB

To deploy and then run the package on Azure SQL Database, you'll need to first **connect** to the SSIS Catalog database (SSISDB). To do so, complete the following steps:

1. Double click **SSMS** icon on desktop to launch **SQL Server Management Studio**
2. In the **Connect to Server** dialog box, enter the following information:
3. **Server name:** azuresql- <
4. **Authentication:** Select **SQL Server Authentication**, as you cannot connect to an Azure SQL Database with Windows authentication
5. **Login:** LabUser
6. **Password:**
7. Click the **Options** button
8. Click on the **Connection Properties** tab
9. Click on the dropdown for selecting a **database**

**NOTE:** If a dialogue box appears and asks to connect first, click **Yes**.

10. Click **SSISDB** and make sure it's highlighted
11. Click **OK**
12. Click **Connect**
13. In the **Object Window** on your left, navigate to **Integration Services Catalogs** and click the **Expand Object** icon (plus sign "+"), then expand the **SSISDB** object as well

### Deploying a Project/Package

14. Right click on **SSISDB**, select **Create Folder**



15. Name the folder **Myproject**, click **ok**
16. Click on the **Projects** node
17. Right-click on the **Projects** node and select **Deploy project** to launch the **Services Deployment Wizard**

**NOTE:** You can deploy a project from the current catalog **or** from the file system.

18. On the **Introduction** page of the wizard, review the introduction, then click **Next** to open the **Select Source** page
19. On the **Select Source** page, select **Project deployment file** and enter **C:\temp\Daily.ETL.ispac**

**NOTE:** To deploy a project that is already deployed to an SSIS catalog database, select **Integration Services** catalog, and then enter the server name and the path to the project in the catalog.

20. On the **Select Destination** page, select

Enter the following information as the **User name** and **Password**:

22. Click **Next**
23. **User name:** **LabUser**
24. **Password:** `<inject key="AzureAdUserPassword"> </inject>`
25. Click **Connect**
26. Click **Next** after you **Connect**
27. On **Validate** page, click **Next**
28. On the **Review** page, review the settings you selected
29. Click **Deploy** to start the deployment process
30. Once complete, the **Results** page will display successes and/or failures of any actions. Once complete, click **Close** to exit the wizard

### *Running a Package*

32. Refresh the **Projects** folder in SSMS
33. Select **Daily ETL**
34. Right-click and select **Execute**
35. After the **Execute Package** dialog box opens, click the **Connection Manages** tab
36. Edit the values for **Password**, **ServerName**, and **Username** on the Destination and Source Databases
37. Click on the Destination DB name **WWI\_DW\_Destination\_DB** to change the settings
38. Change the Password to by clicking the ...on the **Password** field
39. Change the ServerName to by clicking the ... on the **ServerName** field
40. Change the UserName to **Labuser**by clicking the ... on the **UserName** field
41. Click on the source DB name **WWI\_Source\_DB** to change the settings
42. Change the Password to by clicking the ...on the **Password** field

43. Change the ServerName to **database.windows.net** by clicking the ... on the **ServerName** field
44. Change the UserName to **LabUser** clicking the ... on the **UserName** field
45. Click **OK** to run the package
46. Click **yes** to open **Execution Report** to review

## Exercise 4: Power BI

### Overview

I

#### Task 1:

1. Double click **Power BI Desktop** icon on desktop to launch **Power BI Desktop**
2. Upon opening the window, a **Welcome to Power BI Desktop** pop-up will appear. Complete the form as follows:
3. **First Name:** Enter Your First Name
4. **Last Name:** Enter Your Last Name
5. **Email Address:** Enter
6. **Enter your phone number:** Enter 555-555-5555
7. **Country/region:** Select **United States**
8. **Company name:** Enter **Contoso Inc**
9. **Company size:** Select **1000+**
10. **Job Title:** Select **Other**
11. Click **Done**

**Note:** If Power BI Desktop doesn't display a screen showing **Success!** within 30 seconds, you may need to click **Turn on protected mode** on the bottom of your application browser.

12. Exit out of the browser inside the application once it displays **Success!**
13. On the opening screen for Power BI Desktop, click **Get data**,
14. On the **Get Data** blade, click on **SQL Server database**

**Note:** This mimics the connection requirements for our datawarehouse to be able to communicate with Power BI the same way a SQL Server database would.

15. Click **Database**, then **SQL Server database**
16. Click **Connect**
17. Enter for the **Server** box
18. Make sure **Import** is selected under **Data Connectivity mode**
19. Click **OK**
20. On the left-hand side of the window, click the **Database** tab

21. For the **User name** enter LabUser
22. **Password:**
23. Select **DataWarehouse**, Check the box next to **Purchasing.PurchaseOrders**
24. Click **Load**
25. On the left hand column of the screen, click the upper-most of the three icons, which when hovered over with the cursor, should display "Report", click **Report**
26. On the right hand side of the screen, under the **Visualizations** header, click the upper left-hand icon titled **Stacked bar chart**, drag the icon onto the blank canvas in the middle of the screen
27. Under the **Fields** header to the right of the Visualizations menu, click both the **OrderDate** and **SupplierID** checkboxes

### *Stop the Azure-SSIS Integration Runtime*

1. On the desktop open a web browser and navigate to  
<copy><https://portal.azure.com/></copy>
2. Navigate to your Azure Data Factory: **DataFactory**
3. Click the **Author & Monitor** tile to open the Data Factory user interface (UI) on a separate tab
4. On the left hand side navigation bar, click the pen icon (Author)
5. On the Factory Resources page, near the bottom, click the **Connections** button
6. When the Connections tab is displayed, click the **Integration Runtimes** tab
7. Click the Pause/Stop button in the Actions section of the **MyFirstIntegrationRuntime** to stop the service
8. When asked "Are you sure you want to Stop integration runtime", click **Stop**

## Conclusion

In this lab you've learned how to create and configure Integration Runtimes in Azure Data Factory to facilitate cloud-based execution, modification, and monitoring of natively hosted SSIS packages. Then, you've learned how to create a simple yet robust data pipeline to create a reliable data copying strategy. Finally, you learned how the entire process comes together for the end-user by connecting the output data to Power BI Desktop, where you can create and visualize various reports and dashboards to gain insight into your data.